Fairness Training Curves for Amazon Dataset - ERM

ERM

Legend:
- DP-SGD (batch=32, ε=2)
- DP-SGD (batch=32, ε=4)
- DP-SGD (batch=32, ε=8)
- SGD (batch=32)

Axes: Suboptimality (y-axis), Epoch (x-axis)